

```

# Импортируем все необходимые библиотеки

import pandas as pd
import numpy as np
from sqlalchemy import create_engine

# from dotenv import load_dotenv
from catboost import CatBoostClassifier
import os
import gc

# Загружаем переменные окружения из .env файла
load_dotenv()

# Создаем URL для SQLAlchemy
SQLALCHEMY_DATABASE_URL = (
    f"postgresql://"
    f"{os.getenv('POSTGRES_USER')}:{os.getenv('POSTGRES_PASSWORD')}@"
    f"{os.getenv('POSTGRES_HOST')}:{os.getenv('POSTGRES_PORT')}/"
    f"{os.getenv('POSTGRES_DATABASE')}"
)
engine = create_engine(SQLALCHEMY_DATABASE_URL)

### Данные по пользователям с нашими признаками

user_info = pd.read_sql(
    """SELECT * FROM users_info_features_ruslan_prashchurovich""",
    con=engine
)

user_info.head()

```

	user_id	gender	country	city	exp_group	os
source \						
0	200	1	Russia	Degtyarsk	3	Android
ads						
1	201	0	Russia	Abakan	0	Android
ads						
2	202	1	Russia	Smolensk	4	Android
ads						
3	203	0	Russia	Moscow	1	iOS
ads						
4	204	0	Russia	Anzhero-Sudzhensk	3	Android
ads						

	age_category	favorite_topic	user_total_views	user_total_likes	user_ctr
0	young		315	43	0.120112
entertainment					
1	middle_aged		632	58	0.084058
tech					

2	child	541	87	0.138535
politics				
3	young	258	61	0.191223
sport				
4	middle_aged	115	23	0.166667
tech				

	favorite_hour	mean_length_view	mean_length_number_view
mean_length_like \			
0	8	1798.600000	8.133333
1439.400000			
1	22	1636.866667	3.800000
1984.266667			
2	22	449.266667	3.133333
1095.600000			
3	16	1053.800000	2.933333
1898.333333			
4	15	1441.266667	4.333333
1104.066667			

	mean_length_number_like
0	5.200000
1	3.666667
2	3.600000
3	4.800000
4	4.600000

### Посты и топики с нашими признаками

```
posts_info = pd.read_sql(
    """SELECT * FROM posts_info_features_ruslan_prashchurovich""",
    con=engine
)
```

```
posts_info.head()
```

	post_id	text
topic \		
0	1	UK economy facing major risks\n\nThe UK manufa...
business		
1	2	Aids and climate top Davos agenda\n\nClimate c...
business		
2	3	Asian quake hits European shares\n\nShares in ...
business		
3	4	India power shares jump on debut\n\nShares in ...
business		
4	5	Lacroix label bought by US firm\n\nLuxury good...
business		

numbers_count	text_length	word_count	TotalTfIdf	MaxTfIdf
---------------	-------------	------------	------------	----------

```

MeanTfIdf \
0      13      1967      327      8.134297  0.511790
0.000541
1      20      2701      449     11.356141  0.254945
0.000756
2      47      3408      571     10.911138  0.410319
0.000726
3       6      1026      178      6.754817  0.372043
0.000450
4       3       889      155      5.946105  0.526337
0.000396

  TextCluster ... DistanceTo19thCluster DistanceTo20thCluster \
0           8 ...           0.611784           0.581796
1          22 ...           0.587971           0.557299
2           8 ...           0.668490           0.658415
3           8 ...           0.716179           0.683872
4          22 ...           0.554863           0.506563

  DistanceTo21thCluster DistanceTo22thCluster DistanceTo23thCluster
\
0           0.584506           0.587526           0.540042
1           0.577983           0.579828           0.512105
2           0.700094           0.685362           0.642903
3           0.682313           0.680886           0.633085
4           0.484730           0.524628           0.448073

  DistanceTo24thCluster DistanceTo25thCluster total_views_post \
0           0.624294           0.702854           7414
1           0.646764           0.703910           6851
2           0.629109           0.773713           7263
3           0.789145           0.804049           6964
4           0.686375           0.635861           7400

  total_likes_post post_ctr
0           1057  0.142568
1            636  0.092833
2           1116  0.153656
3           1165  0.167289
4           1141  0.154189

[5 rows x 38 columns]

```

### Попробуем обучить, скажем, на 5 миллионах строк (таблица слишком большая: более 70 млн строк)

```
feed_data = pd.read_sql(
    """SELECT * FROM feed_data WHERE action = 'view' LIMIT 5000000""",
    con=engine
)
```

```
feed_data.head()
```

	timestamp	user_id	post_id	action	target
0	2021-12-10 14:55:08	84681	1338	view	1
1	2021-12-10 14:57:33	84681	6282	view	0
2	2021-12-10 14:58:52	84681	4553	view	0
3	2021-12-10 15:01:41	84681	3214	view	0
4	2021-12-10 15:03:50	84681	4360	view	0

### Воспроизведем датафрейм со всеми новыми фичами

```
df = pd.merge(feed_data, posts_info, on="post_id", how="left")
df = pd.merge(df, user_info, on="user_id", how="left")
```

```
df.head()
```

	timestamp	user_id	post_id	action	target	\
0	2021-12-10 14:55:08	84681	1338	view	1	
1	2021-12-10 14:57:33	84681	6282	view	0	
2	2021-12-10 14:58:52	84681	4553	view	0	
3	2021-12-10 15:01:41	84681	3214	view	0	
4	2021-12-10 15:03:50	84681	4360	view	0	

	text	topic
numbers_count \		
0	Iraq advice claim sparks new row\n\nThe Tories...	politics
14		
1	Little Dieter Needs To Fly is another in the r...	movie
4		
2	There is only one film I can think of that mig...	movie
0		
3	#goodnotes this brings home that there are s...	covid
4		
4	If this is the authors and directors idea of a...	movie
0		

	text_length	word_count	...	age_category	user_total_views	\
0	3600	609	...	young	396	
1	2428	445	...	young	396	
2	1086	205	...	young	396	
3	140	23	...	young	396	
4	869	149	...	young	396	

	user_total_likes	user_ctr	favorite_topic	favorite_hour	\
0	65	0.140998	sport	11	
1	65	0.140998	sport	11	

2	65	0.140998	sport	11
3	65	0.140998	sport	11
4	65	0.140998	sport	11

	mean_length_view	mean_length_number_view	mean_length_like	\
0	912.466667	3.066667	1725.466667	
1	912.466667	3.066667	1725.466667	
2	912.466667	3.066667	1725.466667	
3	912.466667	3.066667	1725.466667	
4	912.466667	3.066667	1725.466667	

	mean_length_number_like
0	5.866667
1	5.866667
2	5.866667
3	5.866667
4	5.866667

[5 rows x 58 columns]

### А еще научимся выделять признаки из даты

```
df["hour"] = pd.to_datetime(df["timestamp"]).apply(lambda x: x.hour)
df["month"] = pd.to_datetime(df["timestamp"]).apply(lambda x: x.month)
```

# Циклическое кодирование

```
df["hour_sin"] = np.sin(2 * np.pi * df["hour"] / 24)
df["hour_cos"] = np.cos(2 * np.pi * df["hour"] / 24)
df["month_sin"] = np.sin(2 * np.pi * df["month"] / 12)
df["month_cos"] = np.cos(2 * np.pi * df["month"] / 12)
```

```
df.head()
```

	timestamp	user_id	post_id	action	target	\
0	2021-12-10 14:55:08	84681	1338	view	1	
1	2021-12-10 14:57:33	84681	6282	view	0	
2	2021-12-10 14:58:52	84681	4553	view	0	
3	2021-12-10 15:01:41	84681	3214	view	0	
4	2021-12-10 15:03:50	84681	4360	view	0	

	text	topic
0	Iraq advice claim sparks new row\n\nThe Tories...	politics
14		
1	Little Dieter Needs To Fly is another in the r...	movie
4		
2	There is only one film I can think of that mig...	movie
0		
3	#goodnotes this brings home that there are s...	covid
4		
4	If this is the authors and directors idea of a...	movie

0

	text_length	word_count	...	mean_length_view
mean_length_number_view \				
0	3600	609	...	912.466667
3.066667				
1	2428	445	...	912.466667
3.066667				
2	1086	205	...	912.466667
3.066667				
3	140	23	...	912.466667
3.066667				
4	869	149	...	912.466667
3.066667				

	mean_length_like	mean_length_number_like	hour	month	hour_sin	hour_cos	\
0	1725.466667	5.866667	14	12	-0.500000	-	0.866025
1	1725.466667	5.866667	14	12	-0.500000	-	0.866025
2	1725.466667	5.866667	14	12	-0.500000	-	0.866025
3	1725.466667	5.866667	15	12	-0.707107	-	0.707107
4	1725.466667	5.866667	15	12	-0.707107	-	0.707107

	month_sin	month_cos
0	-2.449294e-16	1.0
1	-2.449294e-16	1.0
2	-2.449294e-16	1.0
3	-2.449294e-16	1.0
4	-2.449294e-16	1.0

[5 rows x 64 columns]

### Уберем все ненужные колонки

```
df = df.drop(
    columns=[
        # 'timestamp', ### timestamp пока оставим
        "action",
        "text",
    ],
    axis=1,
)
```

# Почистим переменные

```

del user_info
del posts_info
del feed_data

gc.collect()

0

### За отсечку возьмем 2021-12-10

X_train = df[df.timestamp < "2021-12-10"].drop(
    columns=["timestamp", "target", "user_id", "post_id"], axis=1
)
X_test = df[df.timestamp >= "2021-12-10"].drop(
    columns=["timestamp", "target", "user_id", "post_id"], axis=1
)

y_train = df[df.timestamp < "2021-12-10"]["target"]
y_test = df[df.timestamp >= "2021-12-10"]["target"]

y_train.shape, y_test.shape
((3875355,), (1124645,))

object_cols = [
    "topic",
    "TextCluster",
    "gender",
    "country",
    "city",
    "exp_group",
    "hour",
    "month",
    "favorite_topic",
    "os",
    "source",
    "age_category",
    "favorite_hour",
]

### Теперь обучим катбуст!

seed = 104773
catboost = CatBoostClassifier(
    iterations=150,
    learning_rate=1,
    depth=5,
    cat_features=object_cols,
    verbose=0,
    random_state=seed,
)

```

```
catboost.fit(X_train, y_train, eval_set=(X_test, y_test), verbose=25)
```

```
0:   learn: 0.3480067 test: 0.3930940 best: 0.3930940 (0)   total:
2.23s remaining: 5m 31s
25:  learn: 0.3372880 test: 0.3823317 best: 0.3823317 (25)  total:
49.1s remaining: 3m 54s
50:  learn: 0.3359490 test: 0.3806361 best: 0.3806361 (50)  total: 1m
35s   remaining: 3m 5s
75:  learn: 0.3352086 test: 0.3803387 best: 0.3803387 (75)  total: 2m
21s   remaining: 2m 18s
100: learn: 0.3345435 test: 0.3795612 best: 0.3795612 (100) total: 3m
7s    remaining: 1m 31s
125: learn: 0.3339486 test: 0.3791217 best: 0.3791173 (121) total: 3m
55s   remaining: 44.8s
149: learn: 0.3335277 test: 0.3790315 best: 0.3789989 (147) total: 4m
40s   remaining: 0us
```

```
bestTest = 0.3789989306
```

```
bestIteration = 147
```

```
Shrink model to first 148 iterations.
```

```
<catboost.core.CatBoostClassifier at 0x7dfd349c0810>
```

```
### Замерим качество работы такой модели
```

```
### Возьмем ROC-AUC
```

```
from sklearn.metrics import roc_auc_score
```

```
print(
    f"Качество на трейне: {roc_auc_score(y_train,
catboost.predict_proba(X_train)[:, 1])}"
)
print(
    f"Качество на тесте: {roc_auc_score(y_test,
catboost.predict_proba(X_test)[:, 1])}"
)
```

```
Качество на трейне: 0.6874009321752084
```

```
Качество на тесте: 0.6657179614328319
```

```
### Из любопытства посмотрим на feature_importance
```

```
import seaborn as sns
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
def plot_feature_importance(importance, names, model_type):
```

```
    # Создадим массивы важности и названий признаков
```



```

feature_importance = np.array(importance)
feature_names = np.array(names)

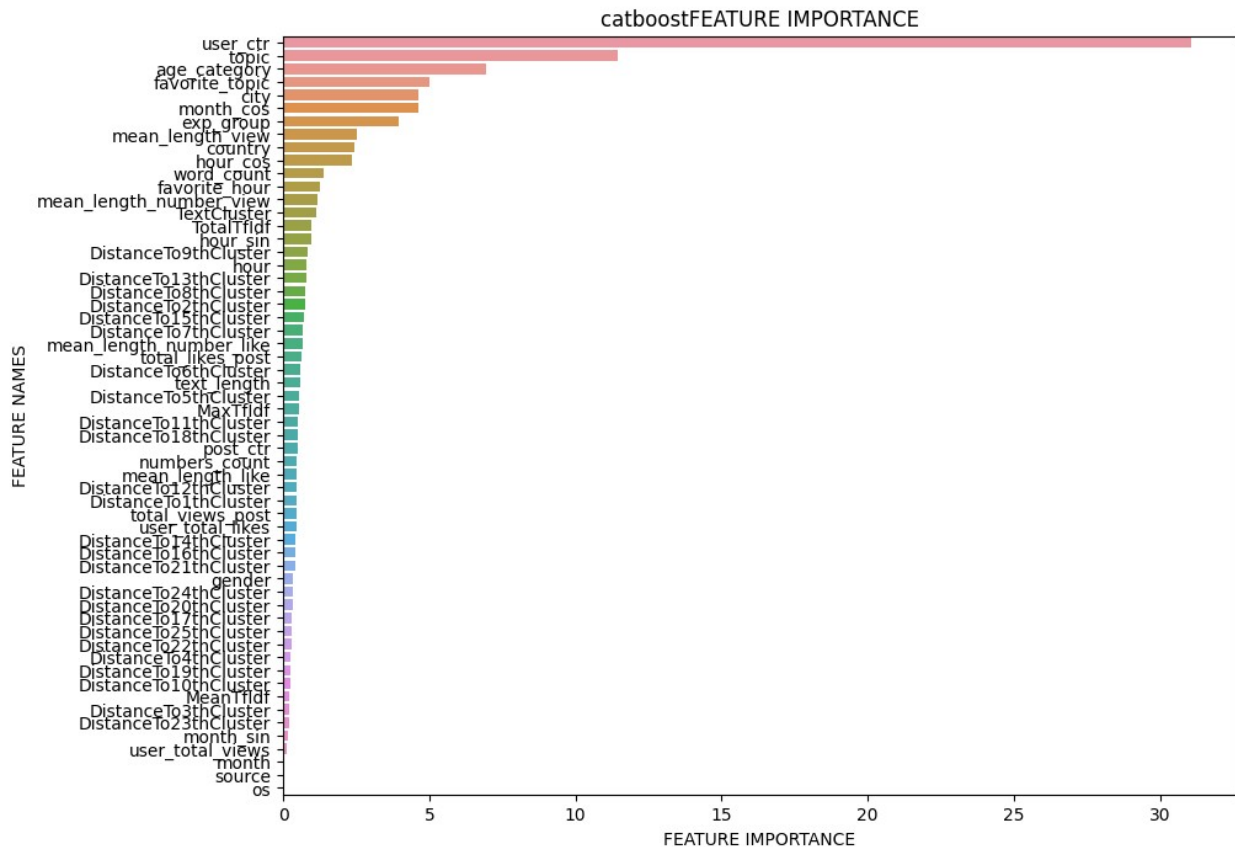
# Создадим датафрейм из словаря
data = {"feature_names": feature_names, "feature_importance":
feature_importance}
fi_df = pd.DataFrame(data)

# Отсортируем по важности
fi_df.sort_values(by=["feature_importance"], ascending=False,
inplace=True)

# Определим размер графика
plt.figure(figsize=(10, 8))
# Нарисуем
sns.barplot(x=fi_df["feature_importance"],
y=fi_df["feature_names"])
# Добавим подписи
plt.title(model_type + "FEATURE IMPORTANCE")
plt.xlabel("FEATURE IMPORTANCE")
plt.ylabel("FEATURE NAMES")

plot_feature_importance(catboost.feature_importances_,
X_train.columns, "catboost")

```



По графику видно, что наиболее важные для модели признаки, полученные в результате feature-engineering, что безусловно, очень круто! Также, стоит отдельно отметить важность признака `age_category`, а значит, решение оставить его было правильным.

### Сохраним модель

```
catboost.save_model("catboost_model", format="cbm")
```